

Introduction to Algorithms

Meng-Tsung Tsai

10/17/2019

Announcements

Written Assignment 2 is due by Oct 31, 15:40. at <https://e3.nctu.me>

Programming Assignment 2 was extended, and is due by Nov 5, 23:59. at <https://oj.nctu.me>

Quiz 1 will be held in class on Oct 24.

Scope: slides 01 - 09, assignments, and their generalizations.

About Quiz 1

Asymptotic Bounds: was1-p1, was1-p2, was2-p1.

Basic DP: was2-p2, pas2-p1.

Reduction: was1-p5, was2-p6.

Don't forget the "I don't know" policy.

You may bring two cheating sheets in A4 size.


Binary Search Trees

What is a binary search tree?

A tree whose keys satisfy the *binary-search-tree property*:

For any node x , all nodes in the left subtree of x has key value at most $x.key$, and all nodes in the right subtree of x has key value at least $x.key$.

--- Example ---



Every node x has 4 attributes:

$x.key$: the key value of x

$x.left$: a pointer to x 's left subtree

$x.right$: a pointer to x 's right subtree

$x.p$: a pointer to x 's parent node


If a child or the parent is missing, the corresponding pointer is set as Null.

In-Order Tree Traversal


The in-order tree traversal visits the nodes in a tree in the order:

for any node x ,

1. the nodes in x 's left subtree are visited before x ,
2. then x ,
3. the nodes in x 's right subtree are visited after x .



In-Order Tree Traversal




The in-order tree traversal visits the nodes in a tree in the order:

for any node x ,

1. the nodes in x 's left subtree are visited before x ,
2. then x ,
3. the nodes in x 's right subtree are visited after x .

In-Order Tree Traversal




The in-order tree traversal visits the nodes in a tree in the order:

for any node x ,

1. the nodes in x 's left subtree are visited before x ,
2. then x ,
3. the nodes in x 's right subtree are visited after x .

In-Order Tree Traversal




The in-order tree traversal visits the nodes in a tree in the order:

for any node x ,

1. the nodes in x 's left subtree are visited before x ,
2. then x ,
3. the nodes in x 's right subtree are visited after x .


Output the keys on a BST in the sorted order



```
In-Order-Traversal(x) {  
    if(x ≠ NIL) {  
        In-Order-Traversal(x.left);  
        print x.key;  
        In-Order-Traversal(x.right);  
    }  
}
```

Q: Why does the above procedure output the keys in the sorted order?


Output the keys on a BST in the sorted order



```
In-Order-Traversal(x) {  
    if(x ≠ NIL) {  
        In-Order-Traversal(x.left);  
        print x.key;  
        In-Order-Traversal(x.right);  
    }  
}
```

Q: Why does the above procedure output the keys in the sorted order?


Output the keys on a BST in the sorted order



```
In-Order-Traversal(x) {  
    if(x ≠ NIL) {  
        In-Order-Traversal(x.left);  
        print x.key;  
        In-Order-Traversal(x.right);  
    }  
}
```

Q: Why does the above procedure output the keys in the sorted order?



Output the keys on a BST in the sorted order



```
In-Order-Traversal(x) {  
    if(x ≠ NIL) {  
        In-Order-Traversal(x.left);  
        print x.key;  
        In-Order-Traversal(x.right);  
    }  
}
```



Q: Why does the above procedure output the keys in the sorted order?

Differences between Heaps and BSTs





Q: Is it possible to output the keys in a heap in the sorted order by an $O(n)$ -time traversal in the comparison-based model?

Differences between Heaps and BSTs





A: No. Otherwise build a heap ($O(n)$ time) followed by the traversal ($O(n)$ time) implements a sorting algorithm. $\rightarrow \leftarrow$

Differences between Heaps and BSTs



Q: Is it possible to build an n -node BST in $o(n \log n)$ time in the comparison-based model?

Differences between Heaps and BSTs



A: No. Otherwise build a BST ($\mathcal{O}(n \log n)$ time) followed by the traversal ($\mathcal{O}(n)$ time) implements a sorting algorithm. $\rightarrow \leftarrow$


Query Operations

Tree-Search(x, k)

Output "Yes" if some node y in the tree rooted at x has value k , or otherwise "No."

--- Example ---

If $x.\text{key} > k$




```
bool Tree-Search(x, k){  
    if(x == Null) return No;  
  
    if(x.key == k) return Yes;  
  
    if(x.key < k)  
        return Tree-Search(x.right, k);  
    else  
        return Tree-Search(x.left, k);  
}
```

Tree-Search(x, k)

Output "Yes" if some node y in the tree rooted at x has value k , or otherwise "No."

--- Running Time ---

If $x.\text{key} > k$




The runtime is $O(\text{height})$.

What is largest possible height of an n -node tree?

Tree-Min(x)

Output the minimum key value in the tree rooted at x has value k, and if no such a value exist, output "-∞."

--- Example ---



```
int Tree-Min(x){  
    if(x == Null) return -∞;  
  
    if(x.left != Null){  
        return Tree-Min(x.left);  
    }else{  
        return x.key;  
    }  
}
```


Exercise

Why is the implementation of Tree-Min correct?


How to implement Tree-Max?

Exercise


Let P be a root-to-leaf path of a binary search tree. Let A be the set of nodes to the left of P . Let B be the set of nodes on P . Let C be the set of nodes to the right of P . Prove or disprove that for any $a \in A$, $b \in B$, $c \in C$, we have $a.\text{key} \leq b.\text{key} \leq c.\text{key}$.



A Counter-Example



Property



$$\max(\ell_1, \ell_2, \dots) \leq B \leq \min(u_1, u_2, \dots)$$


$$\ell_1 \leq \ell_2 \leq \dots \leq u_2 \leq u_1$$

$$A \leq \max(\ell_1, \ell_2, \dots)$$

$$\min(u_1, u_2, \dots) \leq C$$

$$A \leq B \leq C$$

Tree-Successor(x)




Output the smallest key $\geq x.\text{key}$.

(Assume all keys are distinct, or tie-break arbitrarily.)

Tree-Successor(x) is in the tree rooted at x or some u_i . (Why?)

Tree-Predecessor(x)



Output the largest key $\leq x.\text{key}$.

(Assume all keys are distinct, or tie-break arbitrarily.)

Tree-Predecessor(x) is in the tree rooted at x or some ℓ_i . (Why?)

Exercise

Give an implementation of `Tree-successor(x)` so that invoking the function k times can be done in $O(k + \text{height})$ time.

Optimal BST


Optimal BST

Input: n keys $k_1 < k_2 < \dots < k_n$ and a distribution of queried values:

- The queried value is k_i with probability p_i for each i in $[1, n]$.
- The queried value is in-between (k_i, k_{i+1}) (assume it hits a dummy node) with probability q_i for each i in $[1, n]$ where $k_0 = -\infty$ and $k_{n+1} = \infty$.

Output a BST so that the expected search cost is minimized.

--- Example ---



key

k_1

k_2

k_3

d_0

d_1

d_2

d_3

search cost

$2 * p_1$

$1 * p_2$

$2 * p_3$

$3 * q_0$

$3 * q_1$

$3 * q_2$

$3 * q_3$


Optimal BST

Input: n keys $k_1 < k_2 < \dots < k_n$ and a distribution of queried values:

- The queried value is k_i with probability p_i for each i in $[1, n]$.
- The queried value is in-between (k_i, k_{i+1}) (assume it hits a dummy node) with probability q_i for each i in $[1, n]$ where $k_0 = -\infty$ and $k_{n+1} = \infty$.

Output a BST so that the expected search cost is minimized.

--- Example ---



key

k_1

k_2

k_3

d_0

d_1

d_2

d_3

search cost

$1 * p_1$

$3 * p_2$

$2 * p_3$


$2 * q_0$

$4 * q_1$

$4 * q_2$

$3 * q_3$

Optimal BST



Claim. Every t -key subtree has $t+1$ dummy nodes.

count	a key node	a dummy node
parent	1 (except the root)	1
children	2	0


Let X be # dummy node.

By the double counting method on # (parent,child) pairs, we get:

$$t - 1 + X = 2t + 0X,$$

so $X = t+1$.


Optimal BST



Claim. Every subtree that contains keys k_a, k_{a+1}, \dots, k_b also contains dummy nodes d_{a-1}, d_a, \dots, d_b .

Recall the property of BST: $A \leq \max(\ell_1, \ell_2, \dots) \leq B \leq \min(u_1, u_2, \dots) \leq C$,
so the keys and dummy nodes form a continuous interval.

Optimal BST



Let $W(a, b)$ be the sum of probabilities associated with k_a, k_{a+1}, \dots, k_b and d_{a-1}, d_a, \dots, d_b , and $\text{Opt}(a, b)$ be the smallest expected cost of an optimal BST that comprises the same keys and dummy nodes.

Claim. $\text{Opt}(a, b)$

$$= \min_x \text{Opt}(a, x-1) + W(a, x-1) + p_x + \text{Opt}(x+1, b) + W(x+1, b)$$

$$= \min_x \text{Opt}(a, x-1) + \text{Opt}(x+1, b) + W(a, b).$$

Divide and Conquer

OBST(a, b) { // return the smallest expected cost of an optimal BST that comprises keys k_a, k_{a+1}, \dots, k_b and dummy nodes d_{a-1}, d_a, \dots, d_b

if(a-1 == b) return q_b ; // an empty subtree

opt = ∞ ;

for(x = a; x ≤ b; ++x) { // pick a root at k_x

if(OBST(a, x-1)+OBST(x+1, b)+W(a, b) < opt) {

opt = OBST(a, x-1)+OBST(x+1, b)+W(a, b);

}

}

return opt;

}

The initial call is OBST(1, n);

Dynamic Programming

OBST(a, b, sol[][]){ // return the smallest expected cost of an optimal BST that comprises keys k_a, k_{a+1}, \dots, k_b and dummy nodes d_{a-1}, d_a, \dots, d_b

if(a-1 == b) return q_b ; // an empty subtree

if(sol[a][b] < ∞) return sol[a][b];

opt = ∞ ;

for(x = a; x ≤ b; ++x){ // pick a root at k_x

if(OBST(a, x-1, sol)+OBST(x+1, b, sol)+W(a, b) < opt){

opt = OBST(a, x-1, sol)+OBST(x+1, b, sol)+W(a, b);

}

}

return sol[a][b] = opt;

}

The initial call is OBST(1, n, sol = { ∞ });

Output the OBST

OBST(a, b, sol[][]){ // return the smallest expected cost of an optimal BST that comprises keys k_a, k_{a+1}, \dots, k_b and dummy nodes d_{a-1}, d_a, \dots, d_b

if(a-1 == b) return q_b ; // an empty subtree

if(sol[a][b] < ∞) return sol[a][b];

opt = ∞ ;

for(x = a; x ≤ b; ++x){ // pick a root at k_x

if(OBST(a, x-1, sol)+OBST(x+1, b, sol)+W(a, b) < opt){

opt = OBST(a, x-1, sol)+OBST(x+1, b, sol)+W(a, b);

root[a][b] = x; }

}

return sol[a][b] = opt;

}

The initial call is OBST(1, n, sol = { ∞ });

Running Time of OBST:
 $O(n^3)$.

Output the OBST

OutputOBST(a, b) { // output an OBST that comprises keys k_a, k_{a+1}, \dots, k_b
and dummy nodes d_{a-1}, d_a, \dots, d_b (in in-order)

if(a-1 == b) { output d_b ; return; }

OutputOBST(a, root[a][b]-1);

output $k_{\text{root}[a][b]}$;

OutputOBST(root[a][b]+1, b);

}

Exercise

Knuth proves that there always exists an OBST so that

$$\text{root}[i][j-1] \leq \text{root}[i][j] \leq \text{root}[i+1][j] \text{ for every } i, j \text{ in } [1, n].$$

Use this fact to compute OBST **in $O(n^2)$ time**.

Running Time of OBST:
 $O(n^2)$.