

Introduction to Algorithms

Meng-Tsung Tsai

12/03/2019


Dijkstra's algorithm

Problem

Input: a directed graph G , a weight function $\omega : E \rightarrow \mathbf{R}^+ \cup \{0\}$, and a (source) node s in G .

Output: for each node v in G , output the shortest (**may be not simple**) path P from s to v ($\omega(P) = \sum_{e \in P} \omega(e)$ is minimized).

Example.




Problem


Input: a directed graph G , a weight function $\omega : E \rightarrow \mathbf{R}^+ \cup \{0\}$, and a (source) node s in G .

Output: for each node v in G , output the shortest (may be not simple) path P from s to v ($\omega(P) = \sum_{e \in P} \omega(e)$ is minimized).


Example.



A simple case: if $\omega(e)$'s are positive integers




A simple case: if $\omega(e)$'s are positive integers




Because $\omega(e)$'s are 1, running $\text{BFS}(G, s)$ gives you the shortest path to every node in G .

A simple case: if $\omega(e)$'s are positive integers


depth 0



A simple case: if $\omega(e)$'s are positive integers




A simple case: if $\omega(e)$'s are positive integers




A simple case: if $\omega(e)$'s are positive integers


depth 3



A simple case: if $\omega(e)$'s are positive integers



A simple case: if $\omega(e)$'s are positive integers



If some $\omega(e)$ is large, then the corresponding edge will be subdivided into too many nodes, making this approach slow.

Dijkstra's algorithm is an implementation of the BFS approach without subdividing the edges

```
Dijkstra( $G, \omega, s$ ) {  
  Initialization( $G, s$ );  
   $S \leftarrow \emptyset$ ;  
   $Q \leftarrow \{v : v \text{ is a node in } G\}$ ; //  $Q$  is a min-heap, with the  
   $\text{est}(v)$  attributes as keys  
  while( $Q \neq \emptyset$ ) {  
     $u \leftarrow \text{Extract-Min}(Q)$ ; // find  $v$  whose  $\text{est}(v)$  is smallest in  $Q$   
     $S \leftarrow S \cup \{u\}$ ;  
    foreach node  $v$  in  $\text{Adj}[u]$  {  
      Relax( $u, v, \omega$ ); // an implicit Decrease-Key( $v$ )  
    }  
  }  
}
```

Dijkstra's algorithm is an implementation of the BFS approach without subdividing the edges

```
Dijkstra( $G, \omega, s$ ) {  
  Initialization( $G, s$ );  
   $S \leftarrow \emptyset$ ;  
   $Q \leftarrow \{v : v \text{ is a node in } G\}$ ; //  $Q$  is a min-heap, with the  
   $\text{est}(v)$  attributes as keys  
  while( $Q \neq \emptyset$ ) {  
     $u \leftarrow \text{Extract-Min}(Q)$ ; // find  $v$  whose  $\text{est}(v)$  is smallest in  $Q$   
     $S \leftarrow S \cup \{u\}$ ;  
    foreach node  $v$  in  $\text{Adj}[u]$  {  
      Relax( $u, v, \omega$ ); // an implicit Decrease-Key( $v$ )  
    }  
  }
```

Like Prim's algorithm, we cannot find v in a heap efficiently.
We need to use a direct-address table to find v in $O(1)$ time.

Correctness

```
while( $Q \neq \emptyset$ ) {  
     $u \leftarrow \text{Extract-Min}(Q)$ ; // find  $v$  whose  $\text{est}(v)$  is smallest in  $Q$   
     $S \leftarrow S \cup \{u\}$ ;  
    foreach node  $v$  in  $\text{Adj}[u]$  {  
        Relax( $u, v, \omega$ ); // an implicit Decrease-Key( $v$ )  
    }  
}
```

Claim. Let $\underline{\text{est}}(v)$ be the $\text{est}(v)$ at the moment that v is extracted from Q . Let v_1, v_2, \dots, v_n be the order that nodes are extracted from Q . We have $\underline{\text{est}}(v_1) \leq \underline{\text{est}}(v_2) \leq \dots \leq \underline{\text{est}}(v_n)$.

v_1 is s , and $\underline{\text{est}}(s)$ is 0.

v_2 can be relaxed only by $v_1 \Rightarrow$

[1] $\underline{\text{est}}(v_2) = \underline{\text{est}}(v_1) + \omega(v_1, v_2) \geq \underline{\text{est}}(v_1)$, or

[2] $\underline{\text{est}}(v_2) = \infty$. // not relaxed, then $\underline{\text{est}}(v_2) = \dots = \underline{\text{est}}(v_n) = \infty$.

Correctness

```
while( $Q \neq \emptyset$ ) {  
     $u \leftarrow \text{Extract-Min}(Q)$ ; // find  $v$  whose  $\text{est}(v)$  is closest to  $s$   
     $S \leftarrow S \cup \{u\}$ ;  
    foreach node  $v$  in  $\text{Adj}[u]$  {  
        Relax( $u, v, \omega$ ); // an implicit Decrease-Key( $v$ )  
    }  
}
```

Claim. Let $\underline{\text{est}}(v)$ be the $\text{est}(v)$ at the moment that v is extracted from Q . Let v_1, v_2, \dots, v_n be the order that nodes are extracted from Q . We have $\underline{\text{est}}(v_1) \leq \underline{\text{est}}(v_2) \leq \dots \leq \underline{\text{est}}(v_n)$.

v_3 can be relaxed by either v_1 or $v_2 \Rightarrow$

[1] $\underline{\text{est}}(v_3) = \underline{\text{est}}(v_1) + \omega(v_1, v_3) \geq \underline{\text{est}}(v_1)$, or

[2] $\underline{\text{est}}(v_3) = \underline{\text{est}}(v_2) + \omega(v_2, v_3) \geq \underline{\text{est}}(v_2) \geq \underline{\text{est}}(v_1)$, or [3] $\underline{\text{est}}(v_3) = \infty$.

If [3] or $\underline{\text{est}}(v_3) \geq \underline{\text{est}}(v_2)$ holds, then we are done. Otherwise, [1] holds $\Rightarrow \underline{\text{est}}(v_3) = \underline{\text{est}}(v_1) + \omega(v_1, v_3) \geq \underline{\text{est}}(v_2)$ because $\underline{\text{est}}(v_2)$ is the smallest value among $\text{est}(v)$'s relaxed by v_j for some $j < 2$.

Correctness

```
while( $Q \neq \emptyset$ ) {  
     $u \leftarrow \text{Extract-Min}(Q)$ ; // find  $v$  whose  $\text{est}(v)$  is closest to  $s$   
     $S \leftarrow S \cup \{u\}$ ;  
    foreach node  $v$  in  $\text{Adj}[u]$  {  
        Relax( $u, v, \omega$ ); // an implicit Decrease-Key( $v$ )  
    }  
}
```

Claim. Let $\underline{\text{est}}(v)$ be the $\text{est}(v)$ at the moment that v is extracted from Q . Let v_1, v_2, \dots, v_n be the order that nodes are extracted from Q . We have $\underline{\text{est}}(v_1) \leq \underline{\text{est}}(v_2) \leq \dots \leq \underline{\text{est}}(v_n)$.

By repeating this argument over all v_i 's (formally by induction), we complete the proof.

Correctness

```
while( $Q \neq \emptyset$ ) {  
     $u \leftarrow \text{Extract-Min}(Q)$ ; // find  $v$  whose  $\text{est}(v)$  is closest to  $s$   
     $S \leftarrow S \cup \{u\}$ ;  
    foreach node  $v$  in  $\text{Adj}[u]$  {  
        Relax( $u, v, \omega$ ); // an implicit Decrease-Key( $v$ )  
    }  
}
```

Claim. Let $\underline{\text{est}}(v)$ be the $\text{est}(v)$ at the moment that v is extracted from Q . Let v_1, v_2, \dots, v_n be the order that nodes are extracted from Q . We have $\underline{\text{est}}(v_1) \leq \underline{\text{est}}(v_2) \leq \dots \leq \underline{\text{est}}(v_n)$.

Observe that only $\underline{\text{est}}(v_i)$ can be used to relax $\text{est}(v_j)$. If $j < i$, then $\underline{\text{est}}(v_j) \leq \underline{\text{est}}(v_i)$, thus $\text{est}(v_j) = \underline{\text{est}}(v_j)$ after v_j is extracted from Q .

Correctness

```
while( $Q \neq \emptyset$ ) {  
     $u \leftarrow \text{Extract-Min}(Q)$ ; // find  $v$  whose  $\text{est}(v)$  is closest to  $s$   
     $S \leftarrow S \cup \{u\}$ ;  
    foreach node  $v$  in  $\text{Adj}[u]$  {  
        Relax( $u, v, \omega$ ); // an implicit Decrease-Key( $v$ )  
    }  
}
```

Claim. $\text{est}(v_i) = \text{dis}(v_i)$. Let $P = (s, x_1), (x_1, x_2), \dots, (x_k, v_i)$ be a shortest path $s \rightsquigarrow v_i$ (thus $\omega(P) = \text{dis}(v_i)$).

Clearly, this holds for $v_1 = s$. We assume that this holds for v_2, \dots, v_{i-1} . If $\omega(x_k, v_i) > 0$, then $\text{dis}(x_k) < \text{dis}(v_i)$. By the previous claim and the above assumption, $x_k = v_j$ for some $j < i$. Hence,
 $\text{est}(v_i) \leq \text{est}(x_k) + \omega(x_k, v_i) = \text{dis}(v_i)$.

Running time

```
while( $Q \neq \emptyset$ ) {  
     $u \leftarrow \text{Extract-Min}(Q)$ ; // find  $v$  whose  $\text{est}(v)$  is closest to  $s$   
     $S \leftarrow S \cup \{u\}$ ;  
    foreach node  $v$  in  $\text{Adj}[u]$  {  
        Relax( $u, v, \omega$ ); // an implicit Decrease-Key( $v$ )  
    }  
}
```

[1] If we linear scan Q to find the minimum, then it takes $O(n)$ time for each node v . In total, it runs in $O(n^2)$ time. The number of Relax operations is at most $O(m) = O(n^2)$.

[2] If we use binary heaps to find the minimum, then it takes $O(\log n)$ time for each node v . In total, it runs in $O(n \log n)$. Relax() now requires perform Decrease-Key(). Thus, the total running time is $O((n+m) \log n)$.

[3] By Fibonacci-heaps, the total running time is $O(n \log n + m)$.

Exercise

If $\omega(x_k, v_i) = 0$, then x_k might be extracted from Q after v_i .
Please complete this proof.


All-Pairs Shortest Paths

Problem

Input: a directed graph G , a weight function $\omega : E \rightarrow \mathbf{R}$.

Output: for **every pair** of nodes u, v in G , output the shortest (may be not simple) path P from u to v ($\omega(P) = \sum_{e \in P} \omega(e)$ is minimized).

Example.




Problem

Input: a directed graph G , a weight function $\omega : E \rightarrow \mathbf{R}$.

Output: for **every pair** of nodes u, v in G , output the shortest (may be not simple) path P from u to v ($\omega(P) = \sum_{e \in P} \omega(e)$ is minimized).

Example.




Problem

Input: a directed graph G , a weight function $\omega : E \rightarrow \mathbf{R}$.

Output: for **every pair** of nodes u, v in G , output the shortest (may be not simple) path P from u to v ($\omega(P) = \sum_{e \in P} \omega(e)$ is minimized).

Example.




Adjacency matrix representation

Input: a matrix **W** in which (1) $\omega_{ij} = 0$ if $i = j$ (2) $\omega_{ij} = \infty$ if (i, j) is not an edge of the input graph (3) ω_{ij} = the weight of the directed edge (i, j)

Output: a matrix **D** in which d_{ij} denotes the weight of a shortest path $i \rightsquigarrow j$, and a matrix **P** in which p_{ij} denotes the predecessor of j on the shortest path $i \rightsquigarrow j$.

Example.



The nodes are numbered from 1 to n .

$\omega_{12} = 2$, $\omega_{33} = 0$, $\omega_{53} = \infty$.

$d_{12} = 2$, $d_{25} = 2$, $d_{14} = 4$

$p_{14} = 5$, $p_{25} = 3$

The choices of graph representation

graph representation	single-source shortest paths	all-pairs shortest paths
	adjacency list	adjacency matrix

The choices of graph representation

	single-source shortest paths	all-pairs shortest paths
graph representation	adjacency list	adjacency matrix

The adjacency list representation use $O(m)$ space, which may be $o(n^2)$. To store the computed results for all of n^2 pairs, it **already needs** $O(n^2)$ space. Hence, saving space by using adjacency list is pointless for all-pairs shortest paths.

Restriction

There exists **no negative cycle** in the input graph.

Given this restriction, we have:

[1] for every pair of nodes u, v , there exists a **simple** shortest path $u \rightsquigarrow v$. // Otherwise, iteratively cut a cycle on the path until it has no cycle.

[1] \Rightarrow [2] for every pair of nodes u, v , there exists a simple shortest path of **length** at most $n-1$. // We define the length of a path to be the number of edges on the path, and define the weight of a path to be the sum of edge weights on the path.

Solving APSP by SSSP

By SSSP

For each node v , compute the SSSP with the source v . By repeating this **n rounds**, we obtain APSP.

	running time	restriction
Dijkstra's (array)	$O(n (n^2))$	no negative edge
Dijkstra's (binary heaps)	$O(n (n+m) \log n)$	no negative edge
Dijkstra's (Fibonacci heaps)	$O(n (n \log n + m))$	no negative edge
Bellman-Ford	$O(n (nm))$	no negative cycle (from some s ?)

Dynamic Programming

Observation

Because the input graph has no negative cycle, it suffices to find a shortest path $u \rightsquigarrow v$ of length at most $n-1$ for every u, v .

Let the shortest path $u \rightsquigarrow v$ be a shortest path $u \rightsquigarrow k$ plus the edge (k, v) for some **unknown k** . // Why? The subpath $(u \rightsquigarrow k)$ of a shortest path $(u \rightsquigarrow v)$ is also a shortest path.

We can find the unknown k by computing

$$d_{uv} = \ell_{uv}^{(n-1)} = \min_k \ell_{uk}^{(n-2)} + \omega_{kv}$$

where $\ell_{uv}^{(t)}$ denotes the minimum weight of any path from u to v **that have length at most t** .

$$\ell_{uv}^{(1)} = \omega_{uv}$$

By definition, $\ell_{uv}^{(1)}$ is the minimum weight of any path from u to v that have length at most 1. Thus, the path is either **an edge** (u, v) or **a null path** (u, u) .

Given $\ell_{uv}^{(t-1)}$ and ω_{uv} , compute $\ell_{uv}^{(t)}$

Extend-Shortest-Paths($L^{(t-1)}$, W) {

Let $L^{(t)} = \{\infty\}$;

for($u = 1$ to n) {

for($v = 1$ to n) {

for($k = 1$ to n) { // can we exclude both $k = u$ and $k = v$?

if($\ell_{uk}^{(t-1)} + \omega_{kv} < \ell_{uv}^{(t)}$) {

$\ell_{uv}^{(t)} \leftarrow \ell_{uk}^{(t-1)} + \omega_{kv}$;

}

}

}

}

}

Running time and space

Running time:

We need to compute $D = L^{(n-1)}$ from $L^{(1)}$. Since it takes $O(n^3)$ time to obtain $L^{(t)}$ from $L^{(t-1)}$, the total running time is $O(n^4)$.

Space:

While computing $L^{(t)}$ from $L^{(t-1)}$, we no longer need $L^{(t-2)}$. Thus, we need only two L matrices at a time, which uses $O(n^2)$ space.

Matrix P :

The computation of predecessor matrix P is omitted for this algorithm.

Faster Dynamic Programming

Observation

Because the input graph has no negative cycle, it suffices to find a shortest path $u \rightsquigarrow v$ of length at most $n-1$ for every u, v .

Let the shortest path $u \rightsquigarrow v$ be a shortest path $u \rightsquigarrow k$ plus another shortest path (k, v) for some **unknown k** . // Why?
The subpaths of a shortest path $(u \rightsquigarrow v)$ are shortest paths.

We can find the unknown k by computing

$$d_{uv} = \ell_{uv}^{(2^r)} = \min_k \ell_{uk}^{(2^{r-1})} + \ell_{kv}^{(2^{r-1})},$$

where $2^r \geq n-1$.

Note that $L^{(n-1)} = L^{(n)} = L^{(n+1)} = \dots = L^{(\infty)}$.

$$\ell_{uv}^{(1)} = \omega_{uv}$$

By definition, $\ell_{uv}^{(1)}$ is the minimum weight of any path from u to v that have length at most 1. Thus, the path is either **an edge** (u, v) or **a null path** (u, u) .

Given $\ell_{uv}^{(2^{r-1})}$, compute $\ell_{uv}^{(2^r)}$

```
Extend-Shortest-Paths( $L^{(2^{r-1})}$ , W){  
  Let  $L^{(2^r)} = \{\infty\}$ ;  
  for(u = 1 to n){  
    for(v = 1 to n){  
      for(k = 1 to n){  
        if(  $\ell_{uk}^{(2^{r-1})} + \ell_{kv}^{(2^{r-1})} < \ell_{uv}^{(2^r)}$  ){  
           $\ell_{uv}^{(2^r)} \leftarrow \ell_{uk}^{(2^{r-1})} + \ell_{kv}^{(2^{r-1})}$ ;  
        }  
      }  
    }  
  }  
}
```


Running time and space

Running time:

We need to compute $D = L^{(2^r)}$ from $L^{(1)}$, where $2^r \geq n-1$. Since it takes $O(n^3)$ time to obtain $L^{(2^t)}$ from $L^{(2^{t-1})}$, the total running time is $O(n^3 \log n)$.

Space:

By the same argument, the space usage is $O(n^2)$.

Matrix P:

The computation of predecessor matrix P is omitted for this algorithm.

Floyd-Warshall Algorithm

Observation

Let path $u \xrightarrow{k} v$ be the shortest path from u to v and in which each intermediate node, **excluding u and v** , has an identifier $(id) \leq k$.

By definition, we have $\omega(u \rightarrow v) = \omega(u \xrightarrow{n} v)$, and

$$\omega(u \xrightarrow{k} v) = \min(\underbrace{\omega(u \xrightarrow{k-1} v)}_{\text{doesn't pass through node } k}, \underbrace{\omega(u \xrightarrow{k-1} k) + \omega(k \xrightarrow{k-1} v)}_{\text{pass through node } k}).$$

doesn't pass through node k

pass through node k

$$\omega(u \rightsquigarrow^0 v) = \omega(u, v)$$

By definition, $u \rightsquigarrow^0 v$ is the shortest path from u to v and in which each intermediate node, **excluding u and v** , has an identifier $(id) \leq 0$. Since every node has an identifier in $\{1, 2, \dots, n\}$, **$u \rightsquigarrow^0 v$ has no intermediate node.**

Given $\omega(u \rightsquigarrow^{k-1} v)$, compute $\omega(u \rightsquigarrow^k v)$

Floyd-Warshall(W) { // let $F^{(k)}$ be the matrix containing the weight $\omega(u \rightsquigarrow^k v)$ for every u, v

$F^{(0)} = W$;

for($k = 1$ to n) {

$F^{(k)} = F^{(k-1)}$; // doesn't pass through node k

for($u = 1$ to n) {

for($v = 1$ to n) {

if($f_{uk}^{(k-1)} + f_{kv}^{(k-1)} < f_{uv}^{(k)}$) {

$f_{uv}^{(k)} \leftarrow f_{uk}^{(k-1)} + f_{kv}^{(k-1)}$;

}

}}}}

Running time and space

Running time:

There are 3 nested loops and each loop has n iterations, the total running time is thus $O(n^3)$.

Space:

By the same argument, the space usage is $O(n^2)$.

Matrix P:

The computation of predecessor matrix P is shown in the following slides.

Predecessor matrix

We have $P_{uv}^{(0)} = u$ if there exists edge (u, v) and $= \text{NIL}$ otherwise.

If $\omega(u \rightsquigarrow^k v) = \omega(u \rightsquigarrow^{k-1} v)$, then it means the shortest path from u to v doesn't pass through node k . Hence, $P_{uv}^{(k)} = P_{uv}^{(k-1)}$.

If $\omega(u \rightsquigarrow^k v) = \omega(u \rightsquigarrow^{k-1} k) + \omega(k \rightsquigarrow^{k-1} v)$, then the predecessor of v can be found in the latter subpath. That is, $P_{uv}^{(k)} = P_{kv}^{(k-1)}$.

Pseudocode

Floyd-Warshall(W) { // let $F^{(k)}$ be the matrix containing the weight $\omega(u \rightsquigarrow^k v)$ for every u, v

$F^{(0)} = W$; $P_{uv}^{(0)} = (\omega_{uv} < \infty) ? u : \text{NIL}$;

for($k = 1$ to n) {

$F^{(k)} = F^{(k-1)}$; $P^{(k)} = P^{(k-1)}$;

for($u = 1$ to n) {

for($v = 1$ to n) {

if($f_{uk}^{(k-1)} + f_{kv}^{(k-1)} < f_{uv}^{(k)}$) {

$f_{uv}^{(k)} \leftarrow f_{uk}^{(k-1)} + f_{kv}^{(k-1)}$;

$P_{uv}^{(k)} \leftarrow P_{kv}^{(k-1)}$;

}

}}}}

Sparse Graphs

By SSSP

Dijkstra's algorithm needs $O(n^2 \log n + nm)$ time, which is very fast if m is small. However, this approach works only for graphs that has no negative edge.

	running time	restriction
Dijkstra's (array)	$O(n \text{ (} n^2 \text{)})$	no negative edge
Dijkstra's (binary heaps)	$O(n \text{ (} n+m \text{)} \log n)$	no negative edge
Dijkstra's (Fibonacci heaps)	$O(n \text{ (} n \log n + m \text{)})$	no negative edge
Bellman-Ford	$O(n \text{ (} nm \text{)})$	no negative cycle (from some s ?)

For the graphs with negative edges

Idea:

Assign a new weight to each edge so that

[1] for each pair of nodes u, v , every shortest path $u \leadsto v$ before the reweighting remains a shortest path $u \leadsto v$ after the reweighting,

[2] all edges have non-negative weight. (thus, we can use the approach of running Dijkstra's algorithm n times)

Johnson's algorithm - the 1st condition

Given any function $h: V \rightarrow \mathbf{R}$, if we let the new weight function ω'_{uv} be $h(u)-h(v)+\omega_{uv}$, then every path P from u to v has weight

$$\sum_{e \in P} \omega'(e) = h(u)-h(v)+\sum_{e \in P} \omega(e).$$

Since $h(u)-h(v)$ is a constant when we fix u and v , a path is a min ω' -weighted shortest path iff it is a min ω -weighted shortest path.

Johnson's algorithm - the 2nd condition

Create a super source node s , and add edge (s, v) for each node v in the input graph with $\omega_{sv} = 0$.

Running the Bellman-Ford algorithm, we have d_{sv} for each v .

Let $h(u) = d_{su}$ and $h(v) = d_{sv}$. Clearly, $\underbrace{d_{sv}}_{\text{shortest path}} \leq \underbrace{d_{su} + \omega_{uv}}_{\text{a path}}$.

Hence, $\omega'_{uv} = h(u) + \omega_{uv} - h(v) \geq 0$. We are done.

Exercise

Prim's algorithm and Dijkstra's algorithm look similar. What are the differences between these two algorithms?

Exercise

In practice, Floyd-Warshall algorithm is usually implemented as follows:

```
dis[,] =  $\{\infty\}$ ;  
foreach(edge u, v) dis[u, v] =  $\omega[u, v]$ ;  
foreach(nod v) dis[v, v] = 0;  
for(k = 1 to n){  
  for(u = 1 to n){  
    for(v = 1 to n){  
      if(dis[u, k] + dis[k, v] < dis[u, v]){  
        dis[u, v] = dis[u, k] + dis[k, v];  
      }  
    }  
  }  
}
```

How to update the predecessor matrix accordingly?