

Introduction to Algorithms

Meng-Tsung Tsai

12/24/2019

Reminder

23:59, Dec 27

programming
assignment #3

13:30-17:30, Dec 28

programing
quiz #2

Reference

Chapter 8 in "Randomized Algorithms" by Motwani and Raghavan.

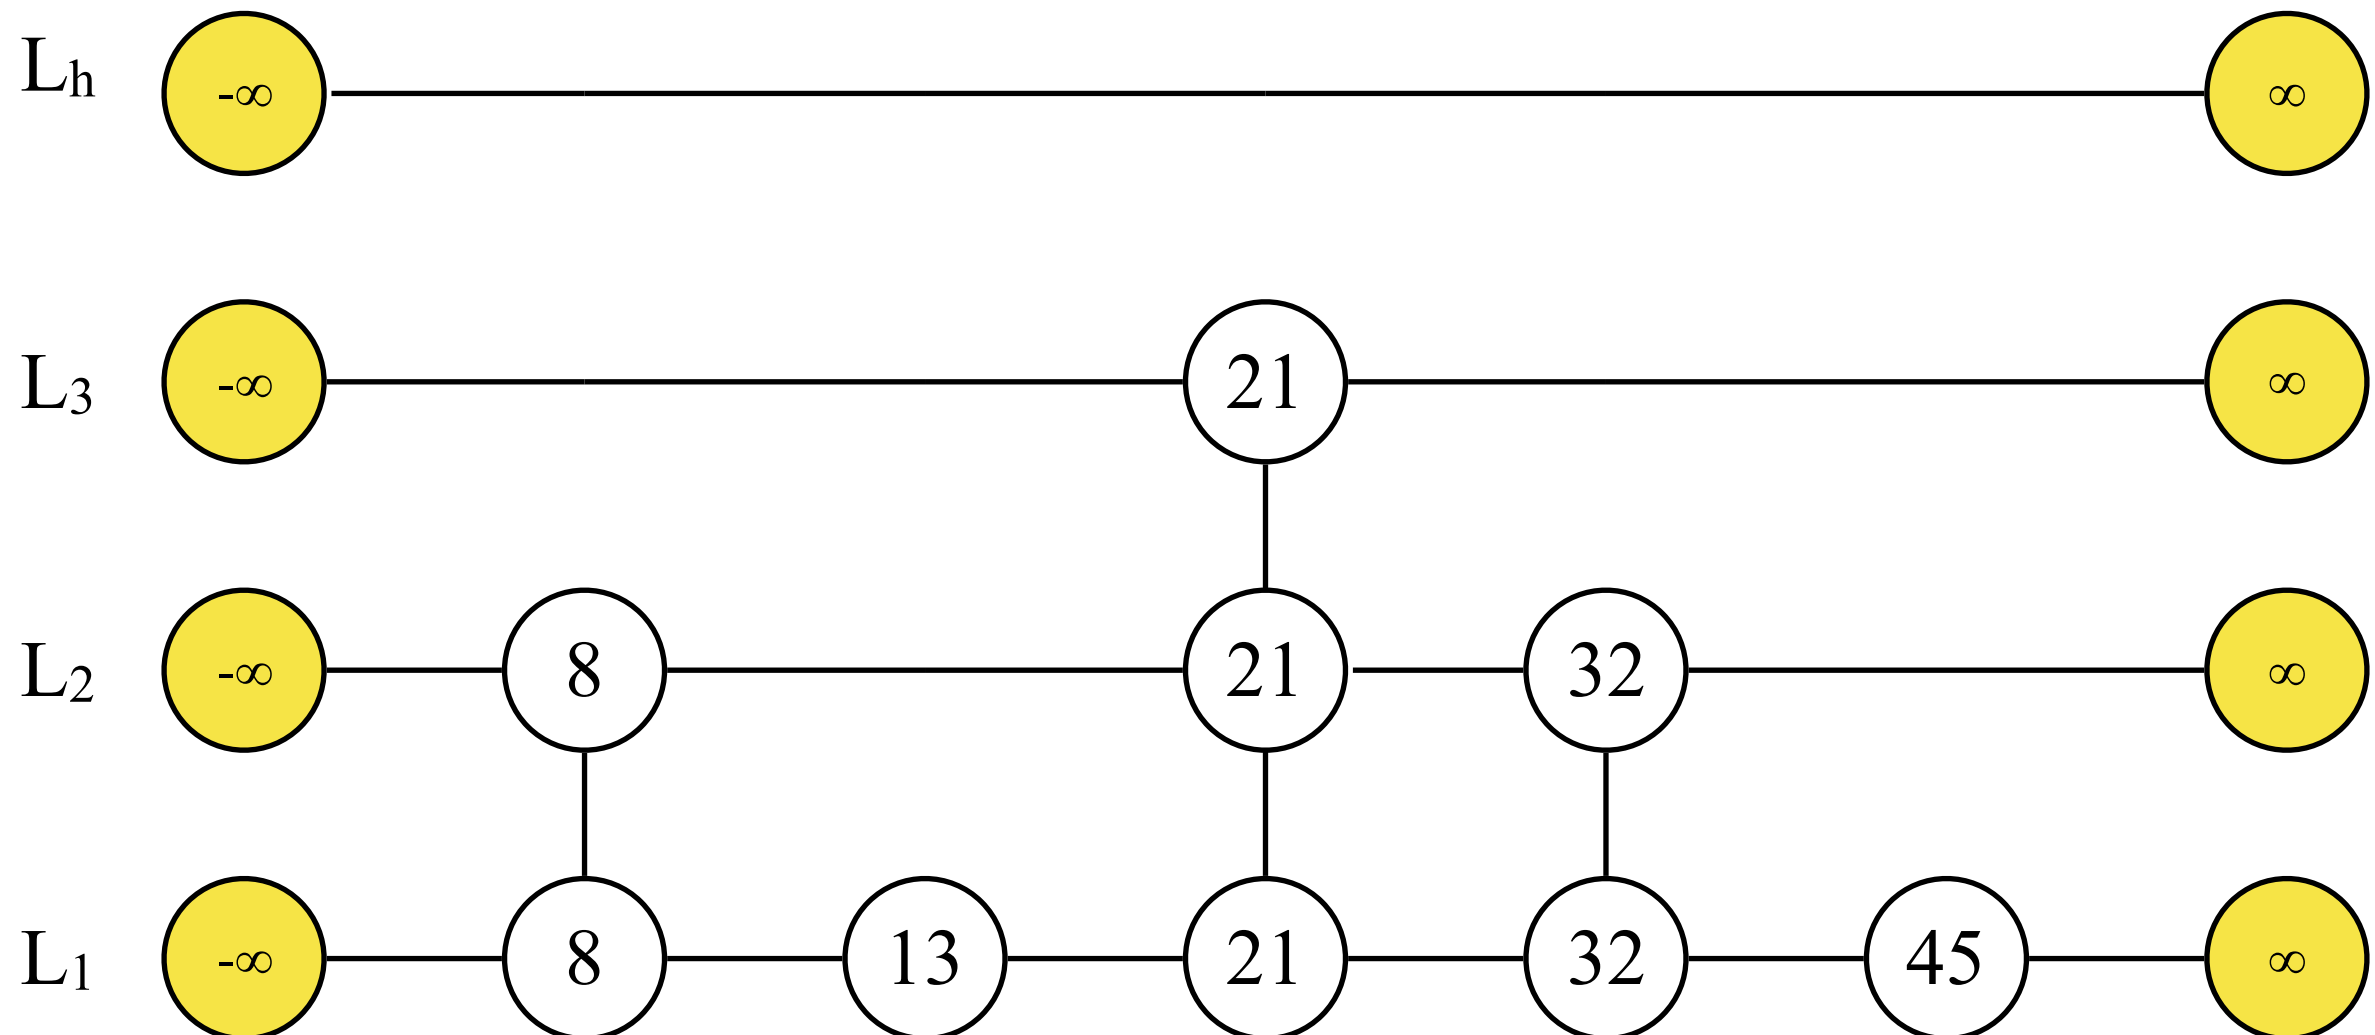
Skip Lists

What is a skip list?

Skip list is a membership container that can support the three operations, (1) insertion, (2) deletion, and (3) search. Each operation takes $O(\log n)$ time on average.

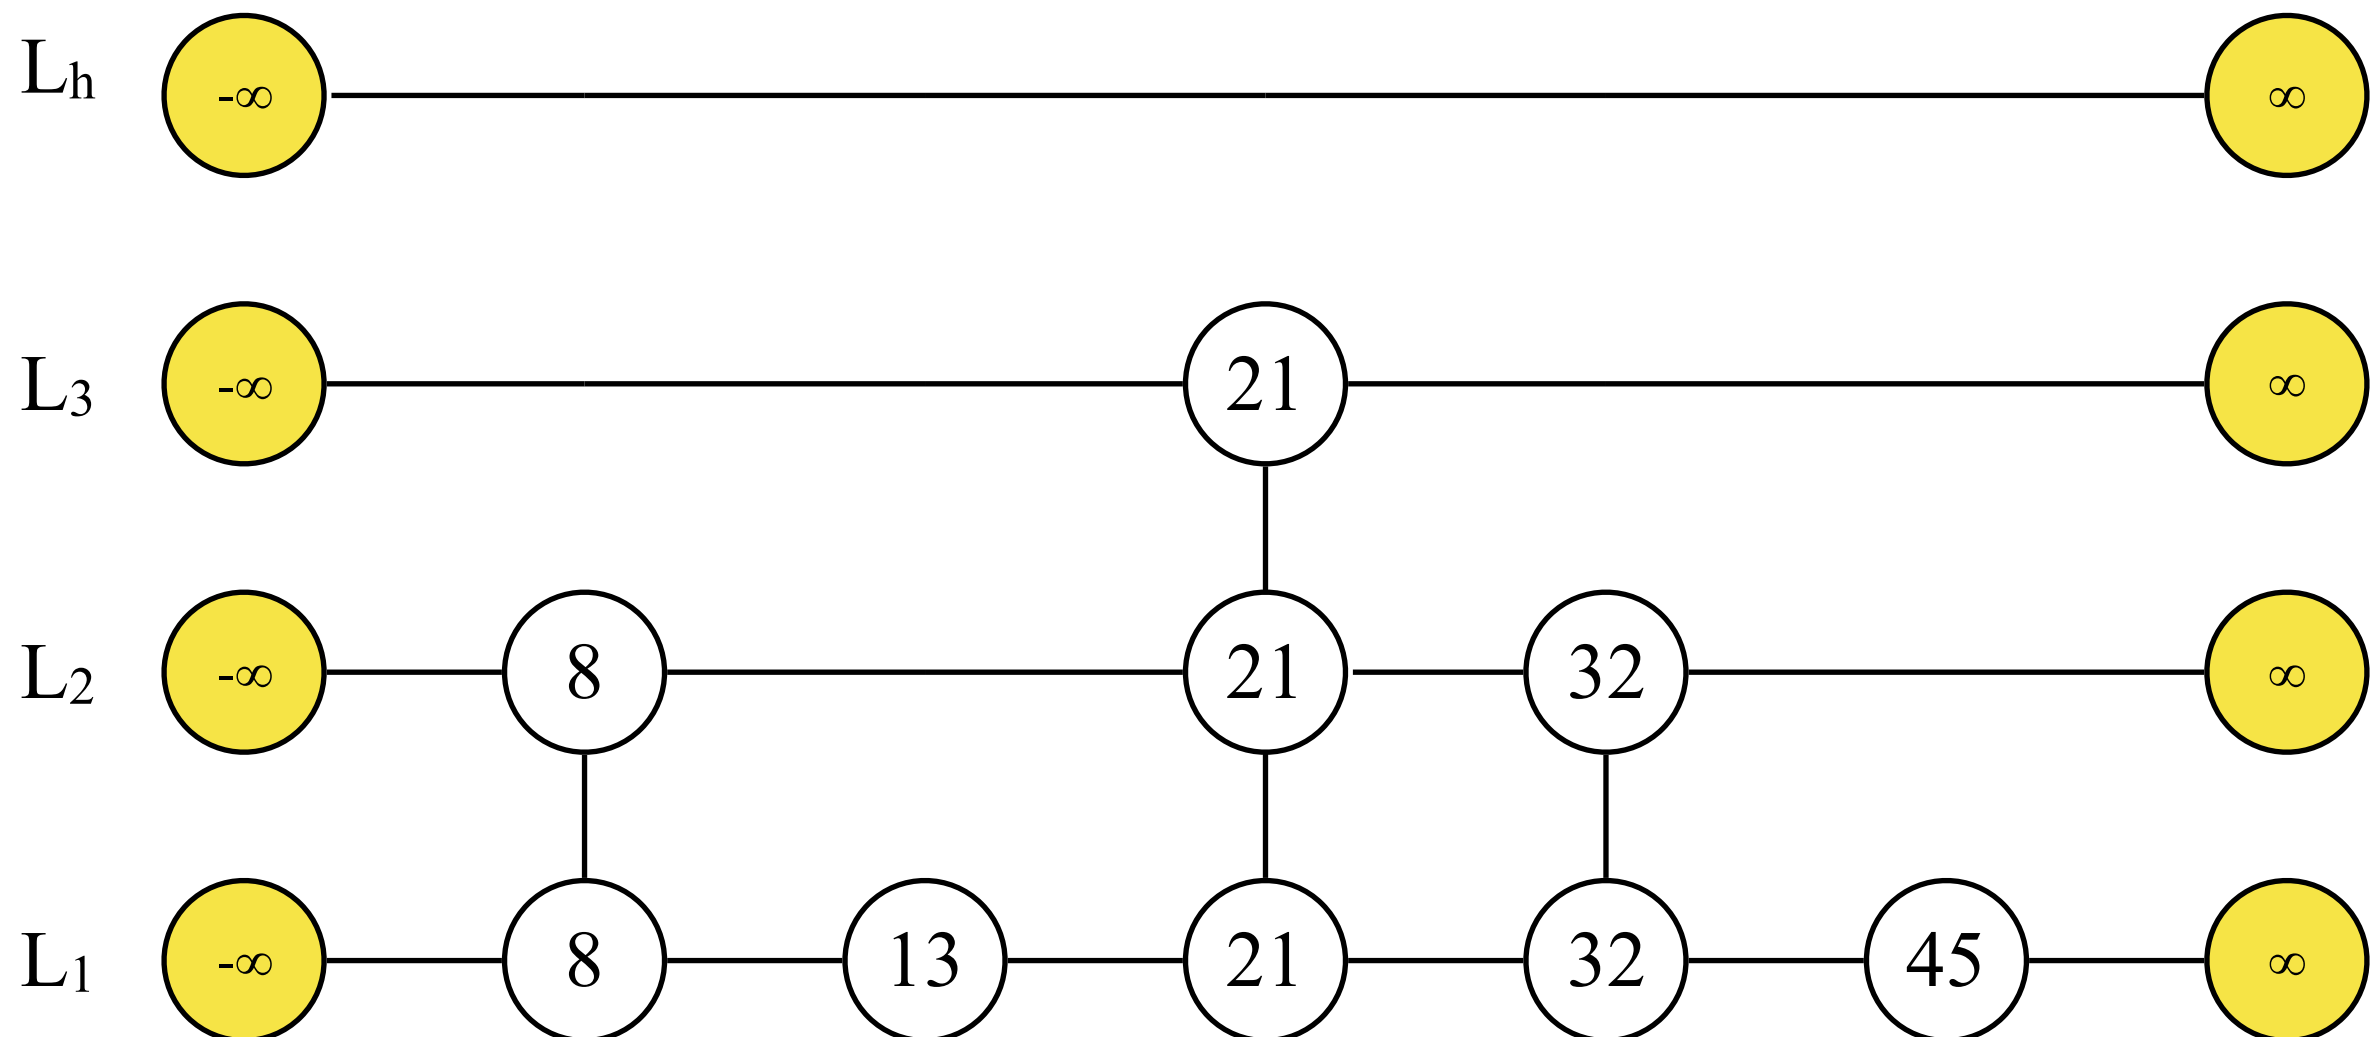
Basics

A skip list is **a collection** of sorted linked list L_1, L_2, \dots, L_h , so that every element in L_i joins L_{i+1} with probability p for every i in $[1, h-1]$ and L_h contains no element. You may assume $p = 1/2$.



Search Operation

Start from L_h and perform a linear scan to find the interval that contains the search key k . Once found, descend to L_{h-1} via vertical pointers and start another linear search in the subinterval. Repeat this while all lists are scanned.



Analysis

The height h is $O(\log n)$ with probability $1 - 1/n^{\Omega(1)}$.

Proof.

For each element, the probability that it joins at least k lists is p^k .

Hence, w.h.p. an element joins $O(\log n)$ lists.

By the Union bound, w.h.p. all elements joins $O(\log n)$ lists.

Analysis

The linear scan performed on a single list takes $O(1)$ steps in expectation.

Proof.

For some list L_i , the probability that the algorithm incurs many probes is small. To see why, every element involves these probes cannot join L_{i+1} . This happens with a tiny probability. The expected length is $O(1)$. Why?

The total search cost is thus $O(\log n)$ because w.h.p. there are $O(\log n)$ lists and each list has expected $O(1)$ probes. It is possible to have more lists, but their total contribution to the expected number of probes is $O(1)$. Why?

Conclusion. The search operation can be done in expected $O(\log n)$ time.

Insertion Operation

The insertion operation works as follows.

Use the search operation to locate the position where the current key shall be placed. Then, insert the key to L_1, L_2, \dots until the random flipped coin tails up.

$O(\log n)$ time for search, and $O(\log n)$ time in expectation to added to the lists.

Conclusion. The insertion operation can be done in $O(\log n)$ time.

Deletion Operation

The deletion operation works as the insertion but in a reverse direction.
The analysis works similarly.

Conclusion. The deletion operation can be done in $O(\log n)$ time.

FKS hashing

What is FKS hashing?

Given n **static** keys, FKS hashing can accommodate these n keys into $O(n)$ space so that each search takes $O(1)$ time in the worst case. The construction time is $O(n)$ on average.

We assume the keys are from the universe $U = \{1, 2, \dots, m\}$ and assume that $m+1$ is a prime p .

FKS hashing is a 2-level hashing. That is, for each search key, we use two hash functions to locate the potential address where the key is placed.

The first level

The first level has **s** slots, and the hash function at the first level is in the following form:

$f_k(x) = ((kx) \bmod p) \bmod s$ where k is sampled uniformly at random from U .

The first level hashing is to (roughly) evenly distribute the keys into s slots.

Analysis

Collision is defined to be the number of pairs of keys in the input that falls within the same slot.

Example. Suppose that $f_k(x)$ hash 5 keys into slot 1 and 1 key into slot 2. Then, the number of collisions is $C(5, 2) + C(1, 2) = 10$.

Let $B(s, r, k, j)$ be the collisions incurred in slot j by using $f_k(x)$ to distribute r keys.

$$\sum_{k=1}^{p-1} \sum_{j=1}^s \binom{B(s, r, k, j)}{2} < (p-1)r^2/s$$

Fix an (x, y) pair, there are at most $2(p-1)/s$ k 's that can hashed them into the same slot. Summing over $C(r, 2)$ pairs, one obtains the above bound.

Analysis

$$\sum_{k=1}^{p-1} \sum_{j=1}^s \binom{B(s, r, k, j)}{2} < (p-1)r^2/s$$

Lemma. By setting $s = r$, there exists some k so that $f_k(x)$ incurs only $O(r)$ collisions.

Corollary. By setting $s = r$, there are a constant fraction of k from U so that $f_k(x)$ incurs only $O(r)$ collisions. --- [Morkov Inequality](#)

Conclusion. After $f_k(x)$ is applied, there are at most $O(r^{1/2})$ keys in a slot.

The second level

The second level has s hash function, each associated to the keys that are hashed into a certain slot. The hash function at the second level is in the following form, again:

$f_k(x) = ((kx) \bmod p) \bmod s$ where k is sampled uniformly at random from U .

The second level hashing is to distribute the keys into s slots without any collision.

Analysis

$$\sum_{k=1}^{p-1} \sum_{j=1}^s \binom{B(s, r, k, j)}{2} < (p-1)r^2/s$$

Lemma. By setting $s = 2r^2$, there exists some k so that $f_k(x)$ incurs **0** collision.

Corollary. By setting $s = 2r^2$, there are a constant fraction of k from U so that $f_k(x)$ incurs 0 collision. --- [Morkov Inequality](#)

Conclusion. After $f_k(x)$ is applied, there are at most 1 key in a slot.

Space Usage

To set $s = n$, one needs $O(n)$ space for the first level hashing.

To set $s = 2r_j^2$, one needs $O\left(\sum_{j=1}^s r_j^2\right)$ space, which is bounded by the number of collision incurred by the first level hashing up to a constant factor. Hence, the space usage of the s hash tables in the second level is $O(n)$ as well.

Probabilistic Verifiers

Matrix Multiplication

Given a pair of n by n matrices A and B , and another matrix C .

Verify whether $AB = C$ using $O(n^2)$ operations with a good prob.

$$\begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & \dots & b_{1n} \\ b_{21} & b_{22} & \dots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \dots & b_{nn} \end{bmatrix}$$

Matrix Multiplication

Given a pair of n by n matrices A and B , and another matrix C .

Verify whether $AB = C$ using $O(n^2)$ operations with a good prob.

$$\begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} \end{bmatrix} \begin{bmatrix} b_{11} & b_{12} & \dots & b_{1n} \\ b_{21} & b_{22} & \dots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{n1} & b_{n2} & \dots & b_{nn} \end{bmatrix}$$

If you multiply A with B directly, then it needs $O(n^3)$ operations.

Key Observation

$$\begin{bmatrix} a_{11} & a_{12} & \dots & a_{1n} \\ a_{21} & a_{22} & \dots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

1. Ax only uses $O(n^2)$ operations.
2. $(AB)x = A(Bx)$
3. If $AB = C$, then for every x , $(AB)x = Cx$.
4. If $AB \neq C$, then for some x , $(AB)x \neq Cx$.

Random Vector \mathbf{x}

$$\begin{bmatrix} p_{11} & p_{12} & \dots & p_{1n} \\ p_{21} & p_{22} & \dots & p_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ p_{n1} & p_{n2} & \dots & p_{nn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

If $\mathbf{AB} \neq \mathbf{C}$, then $\mathbf{P} = \mathbf{AB} - \mathbf{C} \neq \mathbf{0}$, say some entry $p_{ij} \neq 0$.

Sample \mathbf{x} uniformly at random from $\{0, 1\}^n$.

$$\sum_{k=1}^n p_{ik} x_k = p_{ij} x_j + \sum_{k=1, k \neq j} p_{ik} x_k$$

Random Vector x

$$\begin{bmatrix} p_{11} & p_{12} & \dots & p_{1n} \\ p_{21} & p_{22} & \dots & p_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ p_{n1} & p_{n2} & \dots & p_{nn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

If $AB \neq C$, then $\mathbf{P} = \mathbf{AB} - \mathbf{C} \neq \mathbf{0}$, say some entry $p_{ij} \neq 0$.

Sample x uniformly at random from $\{0, 1\}^n$.

$$\sum_{k=1}^n p_{ik} x_k = p_{ij} x_j + \sum_{k=1, k \neq j} p_{ik} x_k$$

With probability $\leq 1/2$, RHS vanishes.

Repetition

Repeat the above procedure multiple (say 3) rounds. Then the verifier succeeds with probability $\geq 1 - (1/2)^3 = 7/8$.

Exercise

Can you use a verifier to implement matrix multiplications?