

Solution to Written Assignment 3

1. Assume w.l.o.g. that $\text{pred}(i) = -1$ initially for every $i \in [0, n-1]$. Define system potential

$$\Phi = \sum_{k \in [0, \sqrt{n}]} n - \text{latest}(k)$$

where $\text{latest}(k)$ is the largest index i that $\text{pred}(i) = k$ if one exists, or 0 otherwise. Let Φ_i be the system potential after i operations are performed. Hence, $\Phi_0 = O(n\sqrt{n})$ and $0 \leq \Phi_n \leq \Phi_0$.

If the i th operation has the actual cost k , then the change of system potential $\Phi_i - \Phi_{i-1} = -k + O(1)$. Thus, the amortized cost $\hat{c}_i = k - k + O(1) = O(1)$. Consequently, the sum of actual cost is upper-bounded by

$$\Phi_0 - \Phi_n + \sum_{i \in [1, n]} \hat{c}_i = O(n\sqrt{n}).$$


2. (a) As shown in Table 1, Fibonacci heaps is the best.

	arrays	binary heaps	Fibonacci heaps
$n \text{ Insert}()$	$O(n)O(1) = O(n)$	$O(n)O(\log n) = O(n \log n)$	$O(n)O(1) = O(n)$
$n \text{ Extract-Min}()$	$O(n)O(n) = O(n^2)$	$O(n)O(\log n) = O(n \log n)$	$O(n)O(\log n) = O(n \log n)$
$m = O(n \log n) \text{ Decrease-Key}()$	$O(m)O(1) = O(n \log n)$	$O(m)O(\log n) = O(n \log^2 n)$	$O(m)O(1) = O(n \log n)$
Total running time	$O(n^2)$	$O(n \log^2 n)$	$O(n \log n)$

Table 1: Running time of different implementations of Prim's algorithm.

- (b) Let \mathcal{H} be the collection of trees. Then, both binary heaps and Fibonacci heaps have the runtime $O(n \log n)$, faster than the runtime of the array implementation $O(n^2)$. Since we assume that binary heaps is easier to implement than Fibonacci heaps, binary heaps is the best for \mathcal{H} .

3. Yes, such a graph G exists, depicted as follows. The roots of BFS and DFS trees are both r .



Solution to Written Assignment 3

4. (a) Replace Line 9 with **for** $k = 1$ to n **do**.
- (b) Algorithm 2 returns an incorrect answer when the n th node is an intermediate node in some shortest path. This bad case can be avoided if the degree of the n th node is 1. Therefore, one may set \mathcal{H} as the collection of all connected graphs whose n th node has degree 1.
5. If such a subsequence exists, then Algorithm 1 returns “Yes” with probability $\geq \delta$ for some constant $\delta > 0$. If such a subsequence does not exist, then Algorithm 1 always returns “No.”

```

1  $x \leftarrow$  an element sampled uniformly at random from  $\{a_1, a_2, \dots, a_n\}$ 
2  $y \leftarrow$  an element sampled uniformly at random from  $\{a_1, a_2, \dots, a_n\} \setminus \{x\}$ 
3  $z \leftarrow$  an element sampled uniformly at random from  $\{a_1, a_2, \dots, a_n\} \setminus \{x\}$ 
4  $\hat{d} \leftarrow \gcd(|x - y|, |x - z|)$ 
5 Use a linear scan to check the intersection between  $a_1, a_2, \dots, a_n$  and
    $\dots, x - 2\hat{d}, x - \hat{d}, x, x + \hat{d}, x + 2\hat{d}, \dots$ 
6 if the intersection contains at least  $n/2$  elements then
7   | output “Yes”
8 else
9   | output “No”
10 end

```

Algorithm 1: Tester.

Algorithm 1 does not always succeed. The undesired events are listed as follows:

- (a) Some of x, y, z is not contained in r_1, r_2, \dots, r_k . This happens with probability

$$1 - \frac{k(k-1)^2}{n(n-1)^2} = 7/8 + o(1).$$

- (b) All of x, y, z are contained in r_1, r_2, \dots, r_k , but $\gcd(p, q) = t$ for some $t > 1$ where $y = x + pd$ and $z = x + qd$. Note that $\hat{d} = d$ if $t = 1$. For a fixed t , this happens with probability at most

$$\frac{k(k-1)^2}{n(n-1)^2} \frac{1}{t^2} = (1/8 - o(1))/t^2.$$

By the Union bound, we have an upper bound on the probability of this undesired event.

$$(1/8 - o(1)) \sum_{t=2}^{\infty} \frac{1}{t^2} < 0.0807$$

Hence, the success rate is at least $1 - 7/8 - 0.0807 - o(1) = \delta > 0$ for some constant $\delta > 0$. If one repeats Algorithm 1 γ times so that

$$(1 - \delta)^\gamma < 1 - 0.99,$$

then the success rate is amplified to at least 0.99, as required. Hence, setting $\gamma = O(1)$ suffices. Since each execution of Algorithm 1 needs $O(n)$ time, the total running time is bounded by $O(\gamma n) = O(n)$.

Solution to Written Assignment 3

6. For $k = 1$, clearly the answer is “Yes.” If $k = 2$, then it suffices to check whether some pair of disks intersect. A simple algorithm can check this in $O(n^2)$ time. Otherwise $k \geq 3$, then such a k -clique, if exists, must have the centers of the corresponding k disks contained in T_{ij} where T_{ij} is defined as the intersection of two disks whose centers are those of d_i and d_j (resp.), and whose radii have length equal the distance r_{ij} between the centers of d_i and d_j , for some $i \neq j \in [1, n]$. Note that one has to ignore those T_{ij} whose $r_{ij} > 2$. For each (i, j) , the disks whose centers are contained in T_{ij} can be found in $O(n)$ time. If there are at least $2k - 4$ centers in the intersection, then a k -clique exists (by the pigeon-hole principle). Otherwise, it suffices to find a largest clique in the subgraph S induced by the disks in the intersection. Note that, in this case, S contains at most $2k - 1 = O(\log n)$ nodes. As mentioned in class, subgraph S is co-bipartite. Finding a largest clique in S can be reduced to finding a maximum matching in the complement graph \bar{S} . By Ford-Fulkerson algorithm, this can be done in $O(k^3) = O(\log^3 n)$ time. Therefore, the total running time is

$$O(n^2(n + \log^3 n)) = O(n^3).$$